Tutorial 6:
AE, VAE, cVAE

Practical Deep Learning for Science
16 May, 2024

מכון ויצמן למדע
WEIZMANN INSTITUTE OF SCIENCE

- Nilotpal

AutoEncoder (AE)

During generation

input image
(x)

Encoder NN

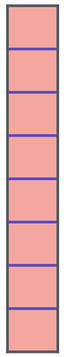

Latent representation
(z)

Decoder NN

output image
$(\bar{x})$

Loss $(x, \bar{x})$

Variational AutoEncoder (VAE)

input image
(x)

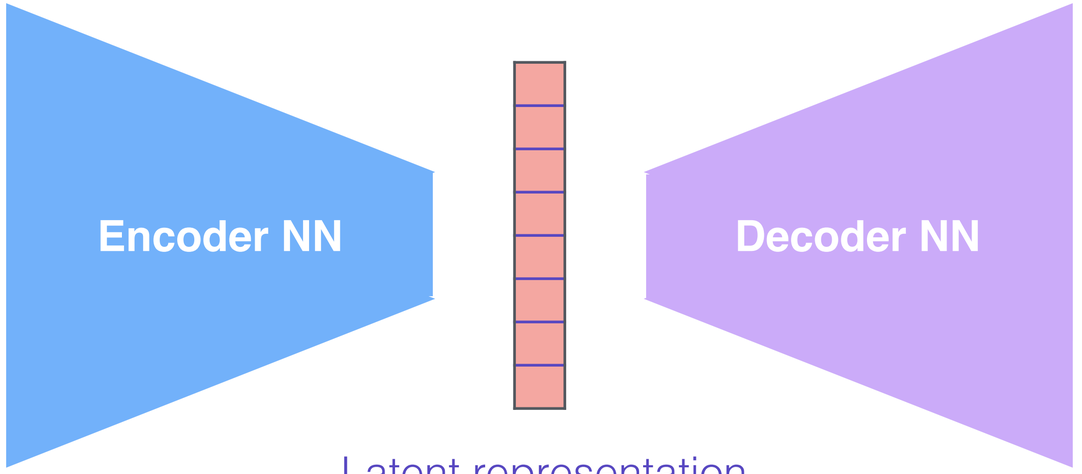

Latent representation
(z)

output image
$(\bar{x})$

Want this to be Normally distributed

$$\text{Loss}\,(x, \bar{x}) + \text{KLD}(z, \text{Gauss})$$

Pointwise KLD

- Target distribution $P = \mathcal{N}(0, I)$
- Predicted distribution $Q = \mathcal{N}(\mu, \Sigma)$
- KLD b/w the two
  - $KLD(P||Q)$
  - `kld_loss = -0.5 * sum(1 + log_var - mu^2 - exp(log_var))`
  - Derivation - https://stats.stackexchange.com/questions/318748/deriving-the-kl-divergence-loss-for-vaes/370048#370048

Conditional Variational AutoEncoder (cVAE)

input image
(x)

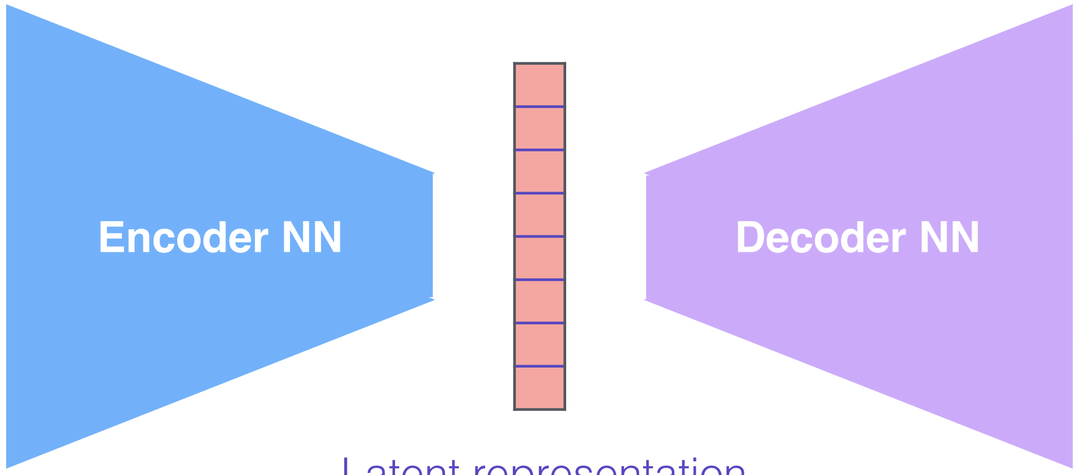

Latent representation
(z)

output image
$(\bar{x})$

- Right now, we have no control over the generated images during inference
  - In the MNIST example, it generates random numbers
  - But let's say we want to generate specific numbers
  - Generation needs to be **conditioned on** what we want → **conditional VAE**

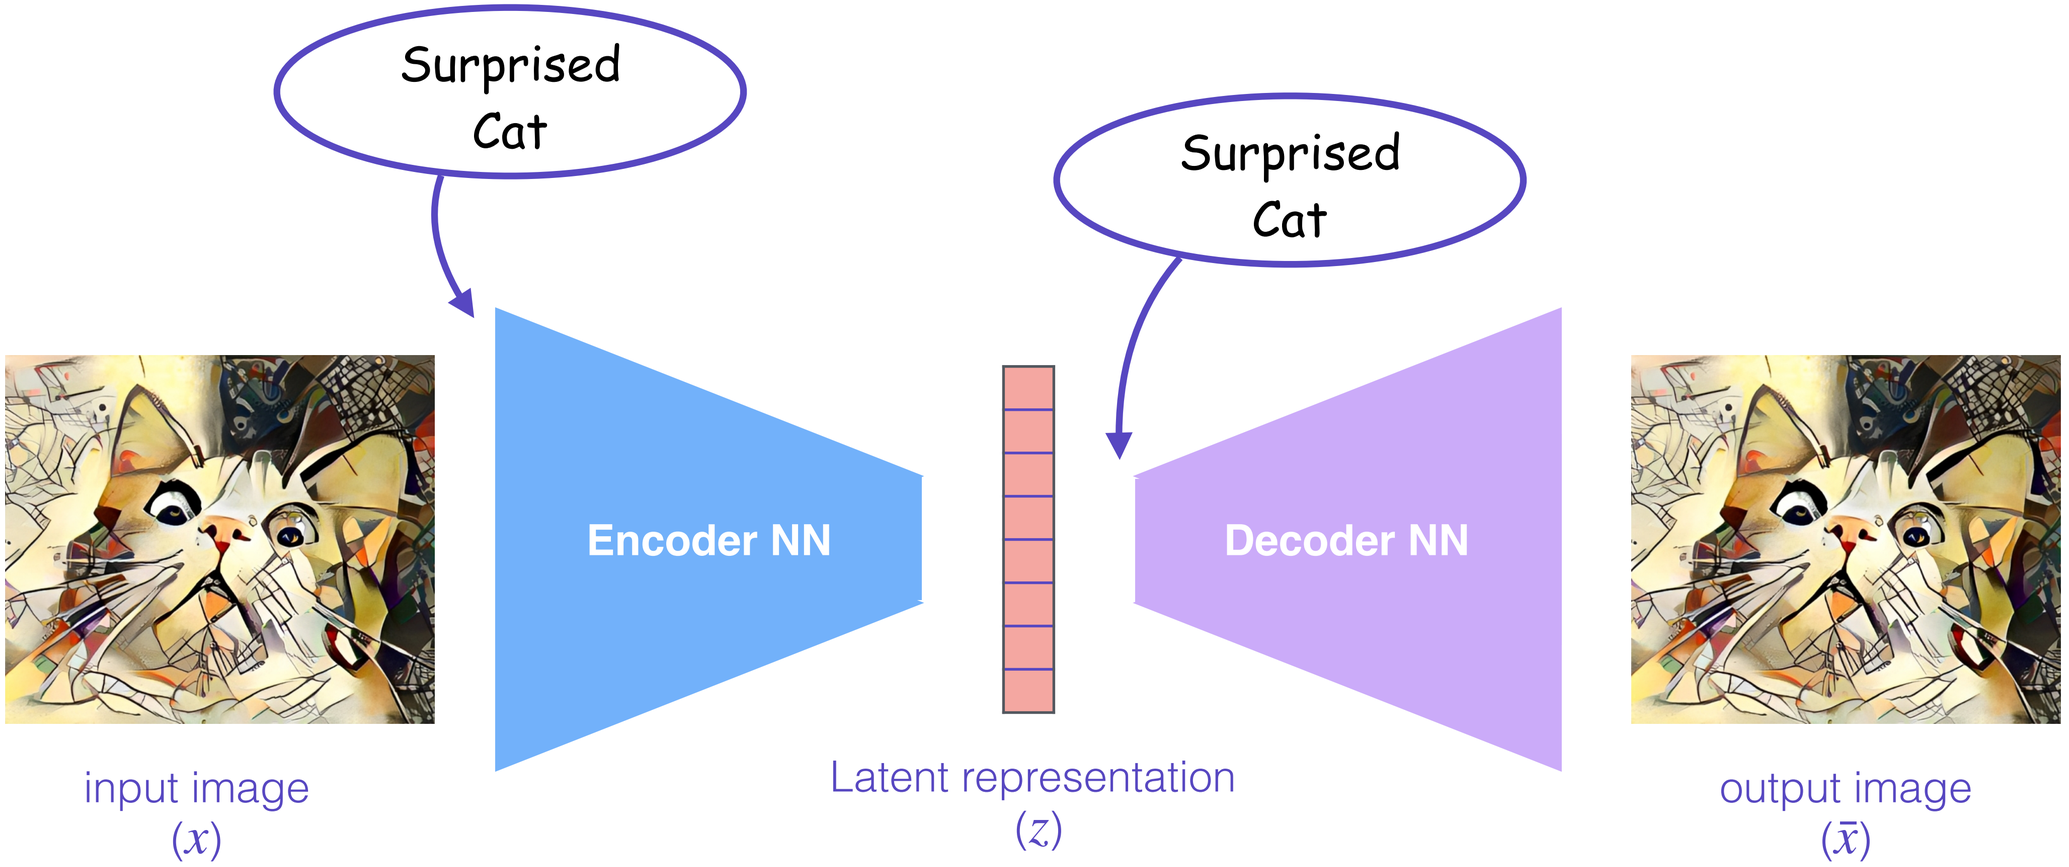

- We need to pass the conditional info as input to both Encoder and Decoder

Encoding the conditional info

- One hot encoded (few classes)
  - 3 → (0, 0, 0, 1, 0, 0, 0, 0, 0, 0)
  - 9 → (0, 0, 0, 0, 0, 0, 0, 0, 0, 1)
  - …

- Encoding class labels into a vector space,
  - `Torch.nn.Embedding()`
  - Helpful when we have a large number of classes

- Other Fancy encoding in text to image models (MIdjourney, DALLE etc)
  - Maybe later in the course (no promises!)

The latent space distributions

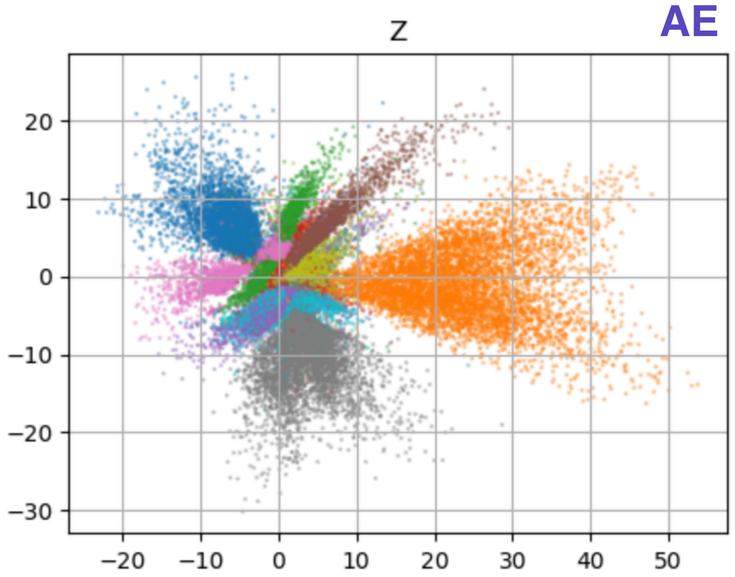

- No KLD → not Gaussian at all
- Empty spaces in between
  - ➡ Sampling is tricky

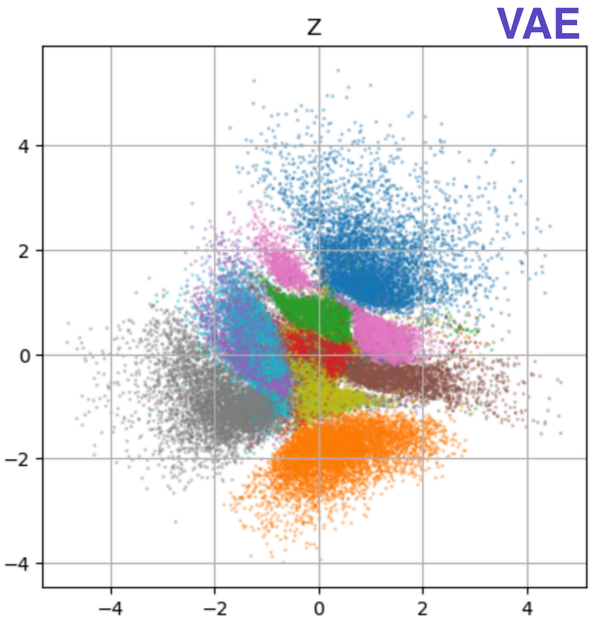

- KLD → Gaussian like
- Less empty spaces in between
  - ➡ Easy to sample
- Needs to keep the classes separate
  - ➡ A 7 must be reconstructed as a 7

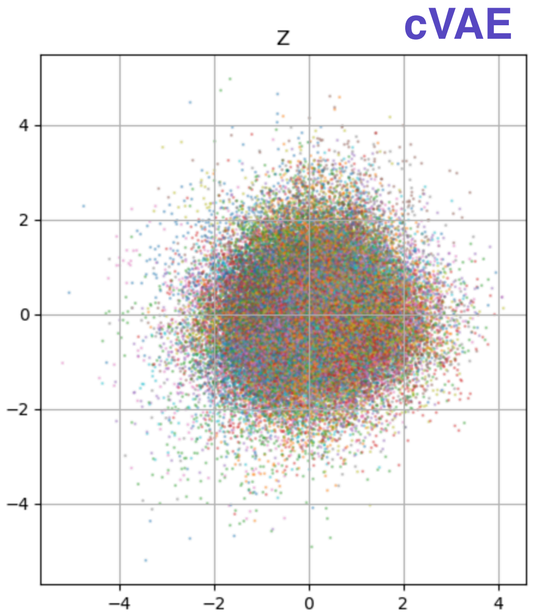

- KLD + conditional info → almost perfect Gaussian
- Doesn't need to keep the classes separate
  - ➡ Already knows which class to reconstruct from the conditional info