Tutorial 7: DCGAN

Practical Deep Learning for Science
23 May, 2024

מכון ויצמן למדע
WEIZMANN INSTITUTE OF SCIENCE

- Nilotpal

GAN

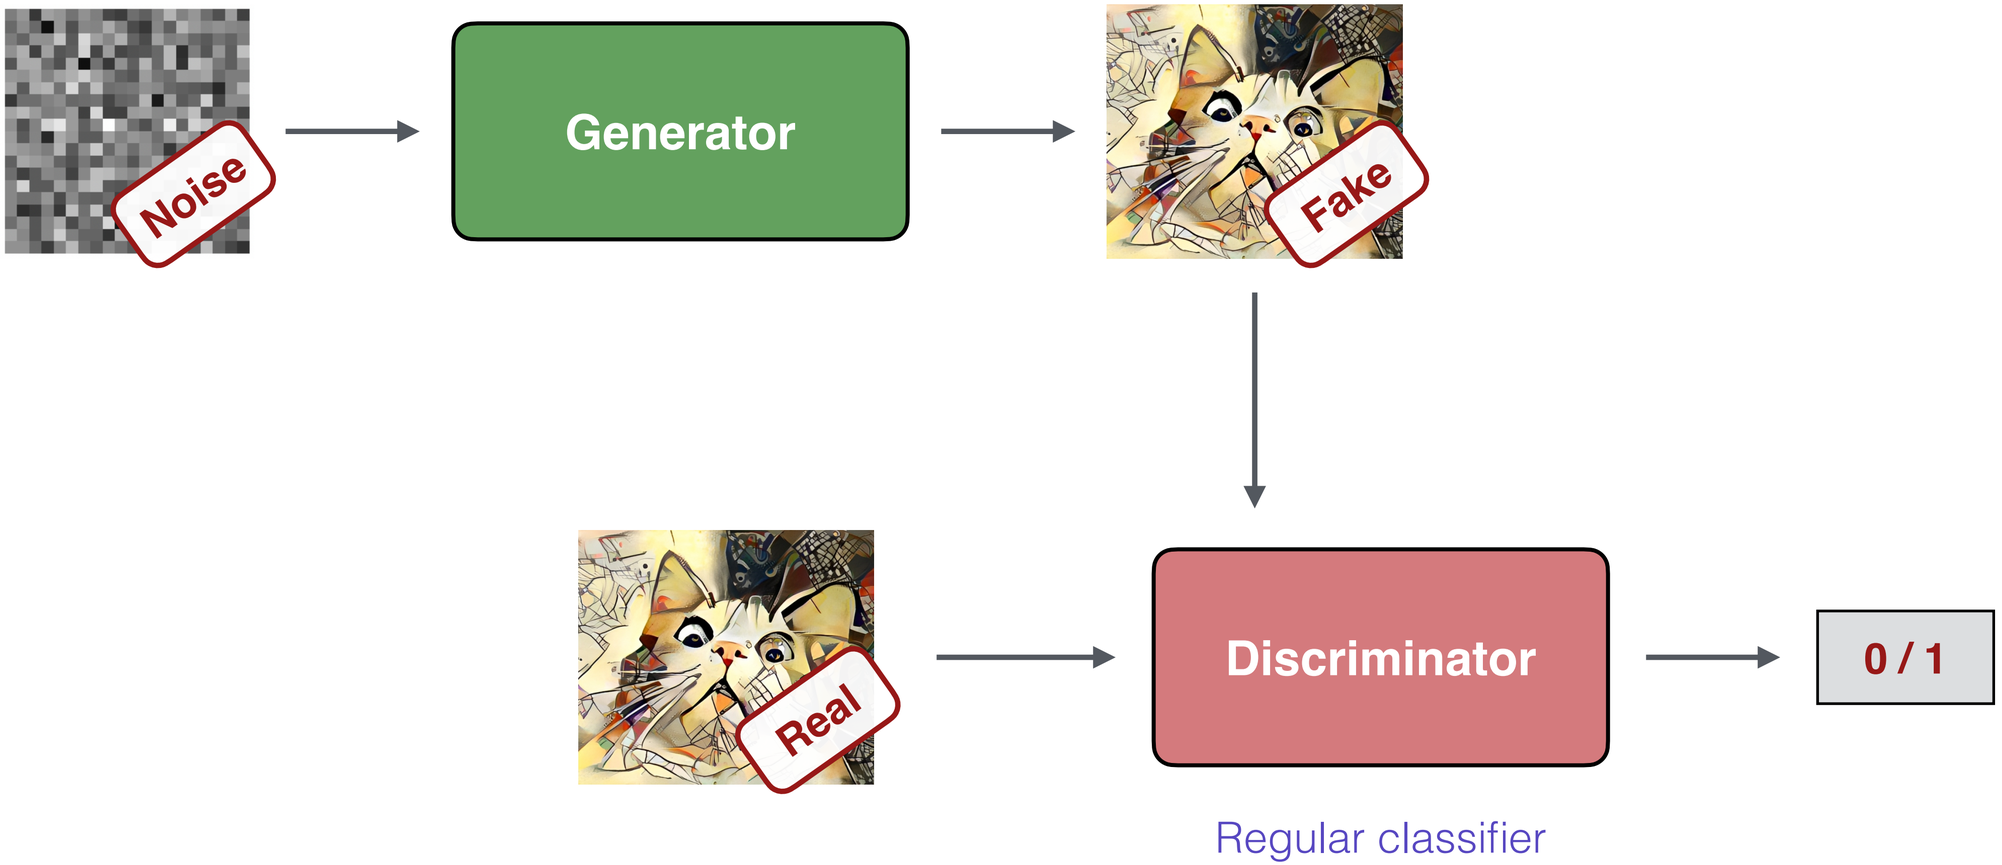
Noise
Generator
Fake
Real
Discriminator
0 / 1
Regular classifier

Loss? It’s a bit tricky

Discriminator is a regular binary classifier

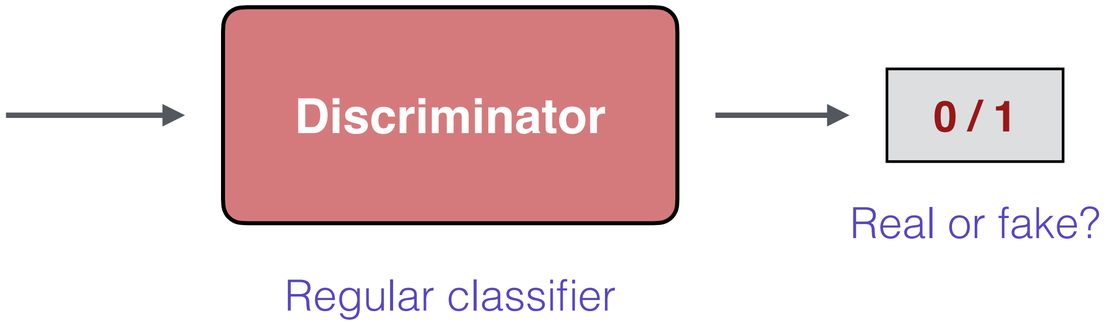

Regular classifier

Binary Cross entropy loss

Generator loss?

input image
(x)

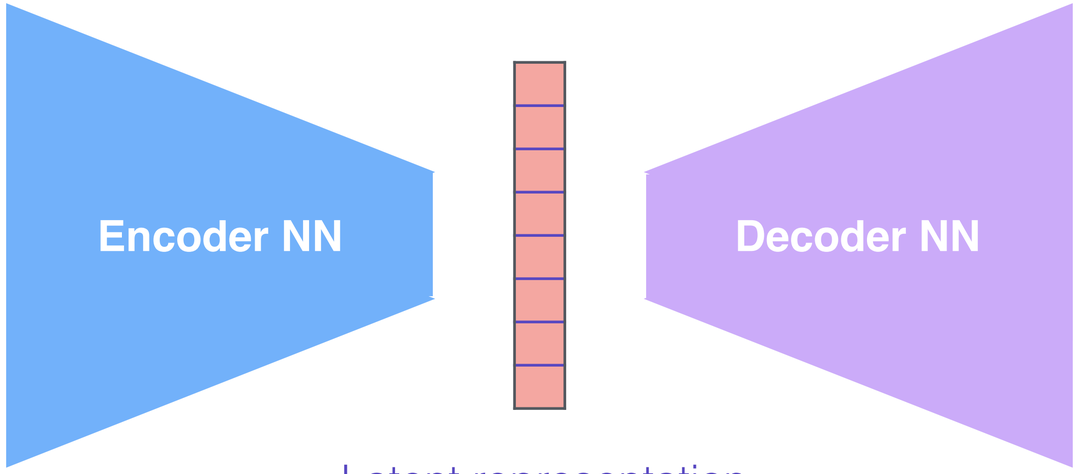

Latent representation
(z)
Noise

output image
$(\bar{x})$

AutoEncoder

Loss (input image, output image

Generator loss?

input image
(x)

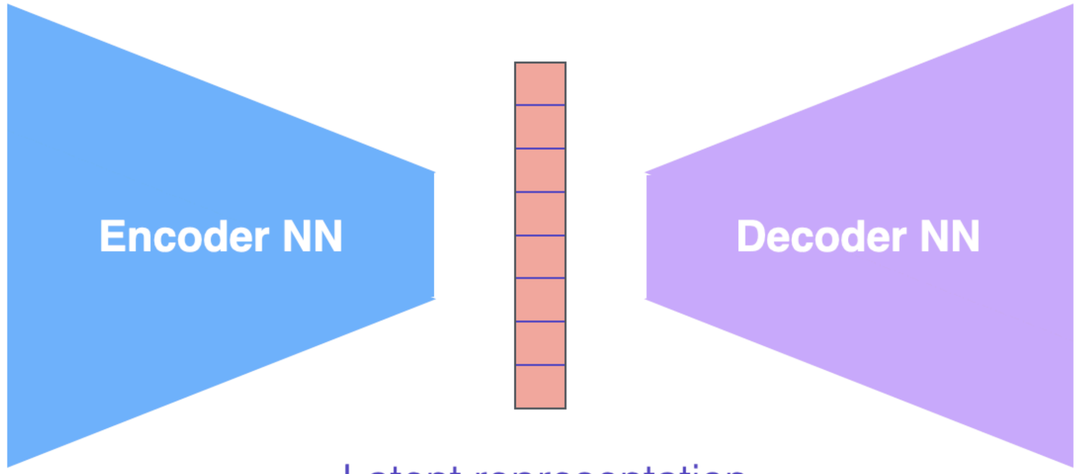

Latent representation
(z)
Noise

output image
$(\bar{x})$

AutoEncoder

Loss (input image, output image

Generator loss?

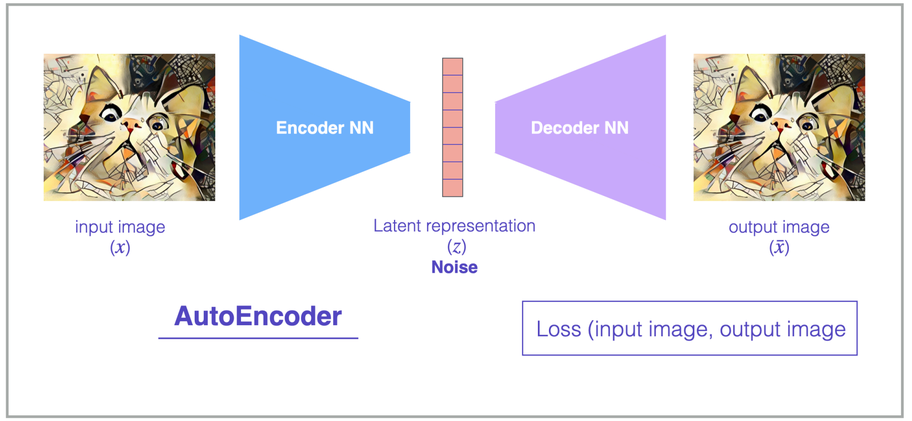

No input image
Can't compute loss like in AE

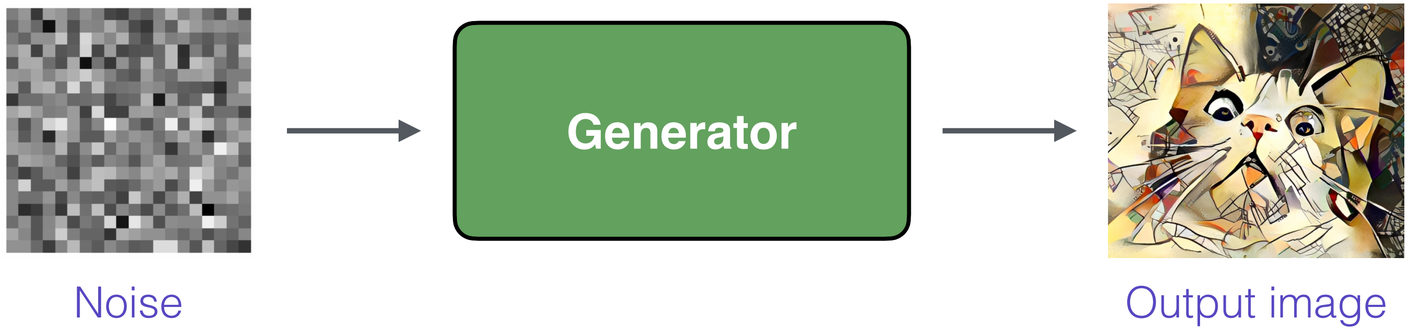

Adversary

- Generator and Discriminator have exact opposite goal
- Will use the discriminator loss again
  - But we will flip the label. *Label the generated images as 1 (real)*
  - Want to increase the likelihood of the fake image being classified as real

- From the discriminator perspective
  - Classify real images as real, fake images as fake
- From Generator's perspective
  - *Generate fake images that gets classified as real*

GAN recipe

- Generate fake images
- Take real images
- Compute discriminator loss; update discriminator

- Label the fake images as real
- Compute discriminator loss; update generator

.detach()

During **discriminator training,** Torch will see **(Generator + Discriminator)** as **one big network**

Losses will get back propagated to generator as well

We need to detach the image from the generator with `fake_img.detach()`

Transposed Convolution

Cons vs TransConv

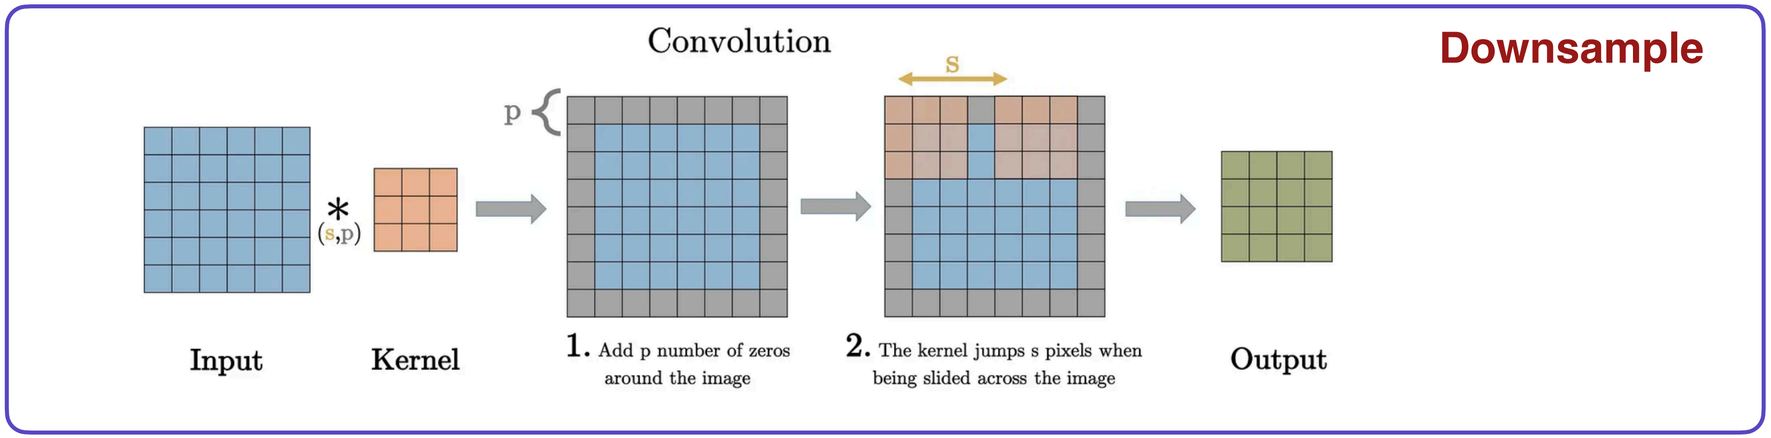

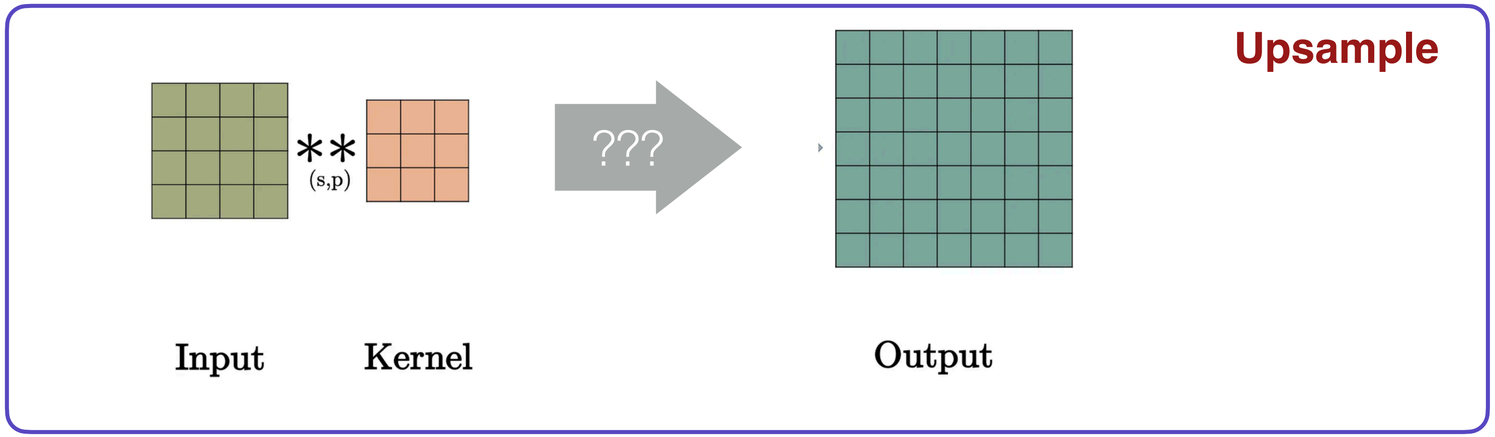

TransConv

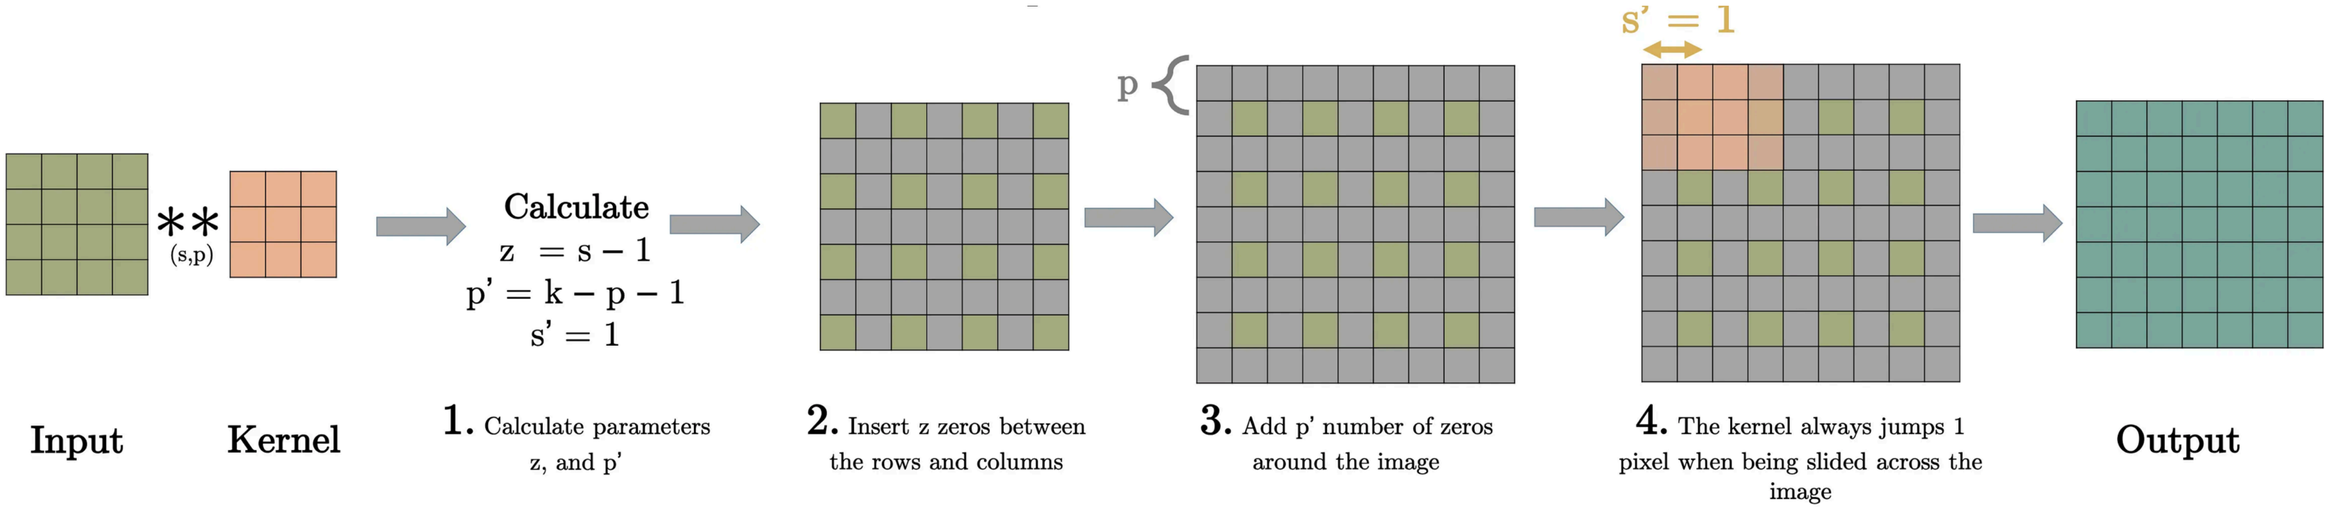

Checkout this article - https://towardsdatascience.com/what-is-transposed-convolutional-layer-40e5e6e31c11
(The graphics are taken from there)

Type: transposed conv - Stride: 1 Padding: 0

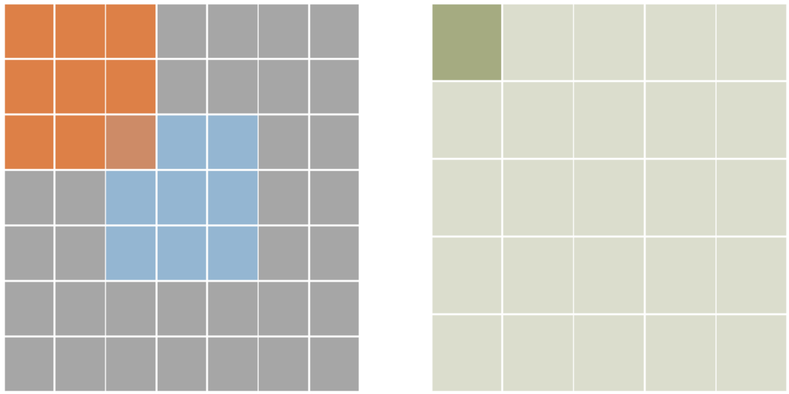

Type: transposed conv - Stride: 1 Padding: 1

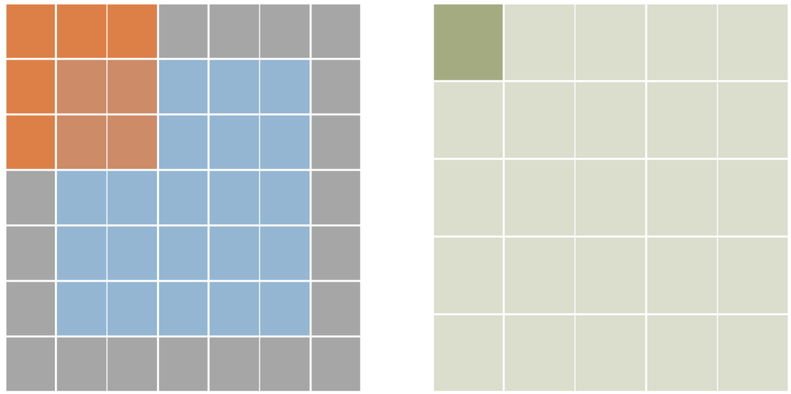

Type: transposed conv - Stride: 2 Padding: 0

Type: transposed conv - Stride: 2 Padding: 1

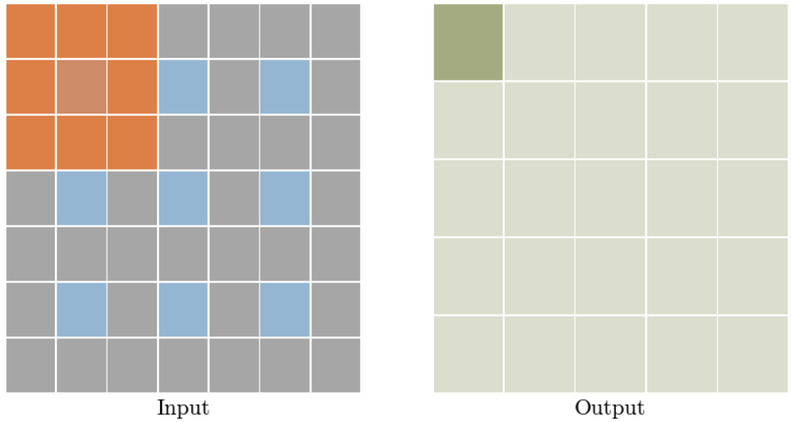

Compute

$$z = s - 1$$
$$p' = k - p - 1$$
$$s' = 1$$

Output size

$$o = (i - 1) \times s + k - 2p$$

k = kernel
p = padding
s = stride