

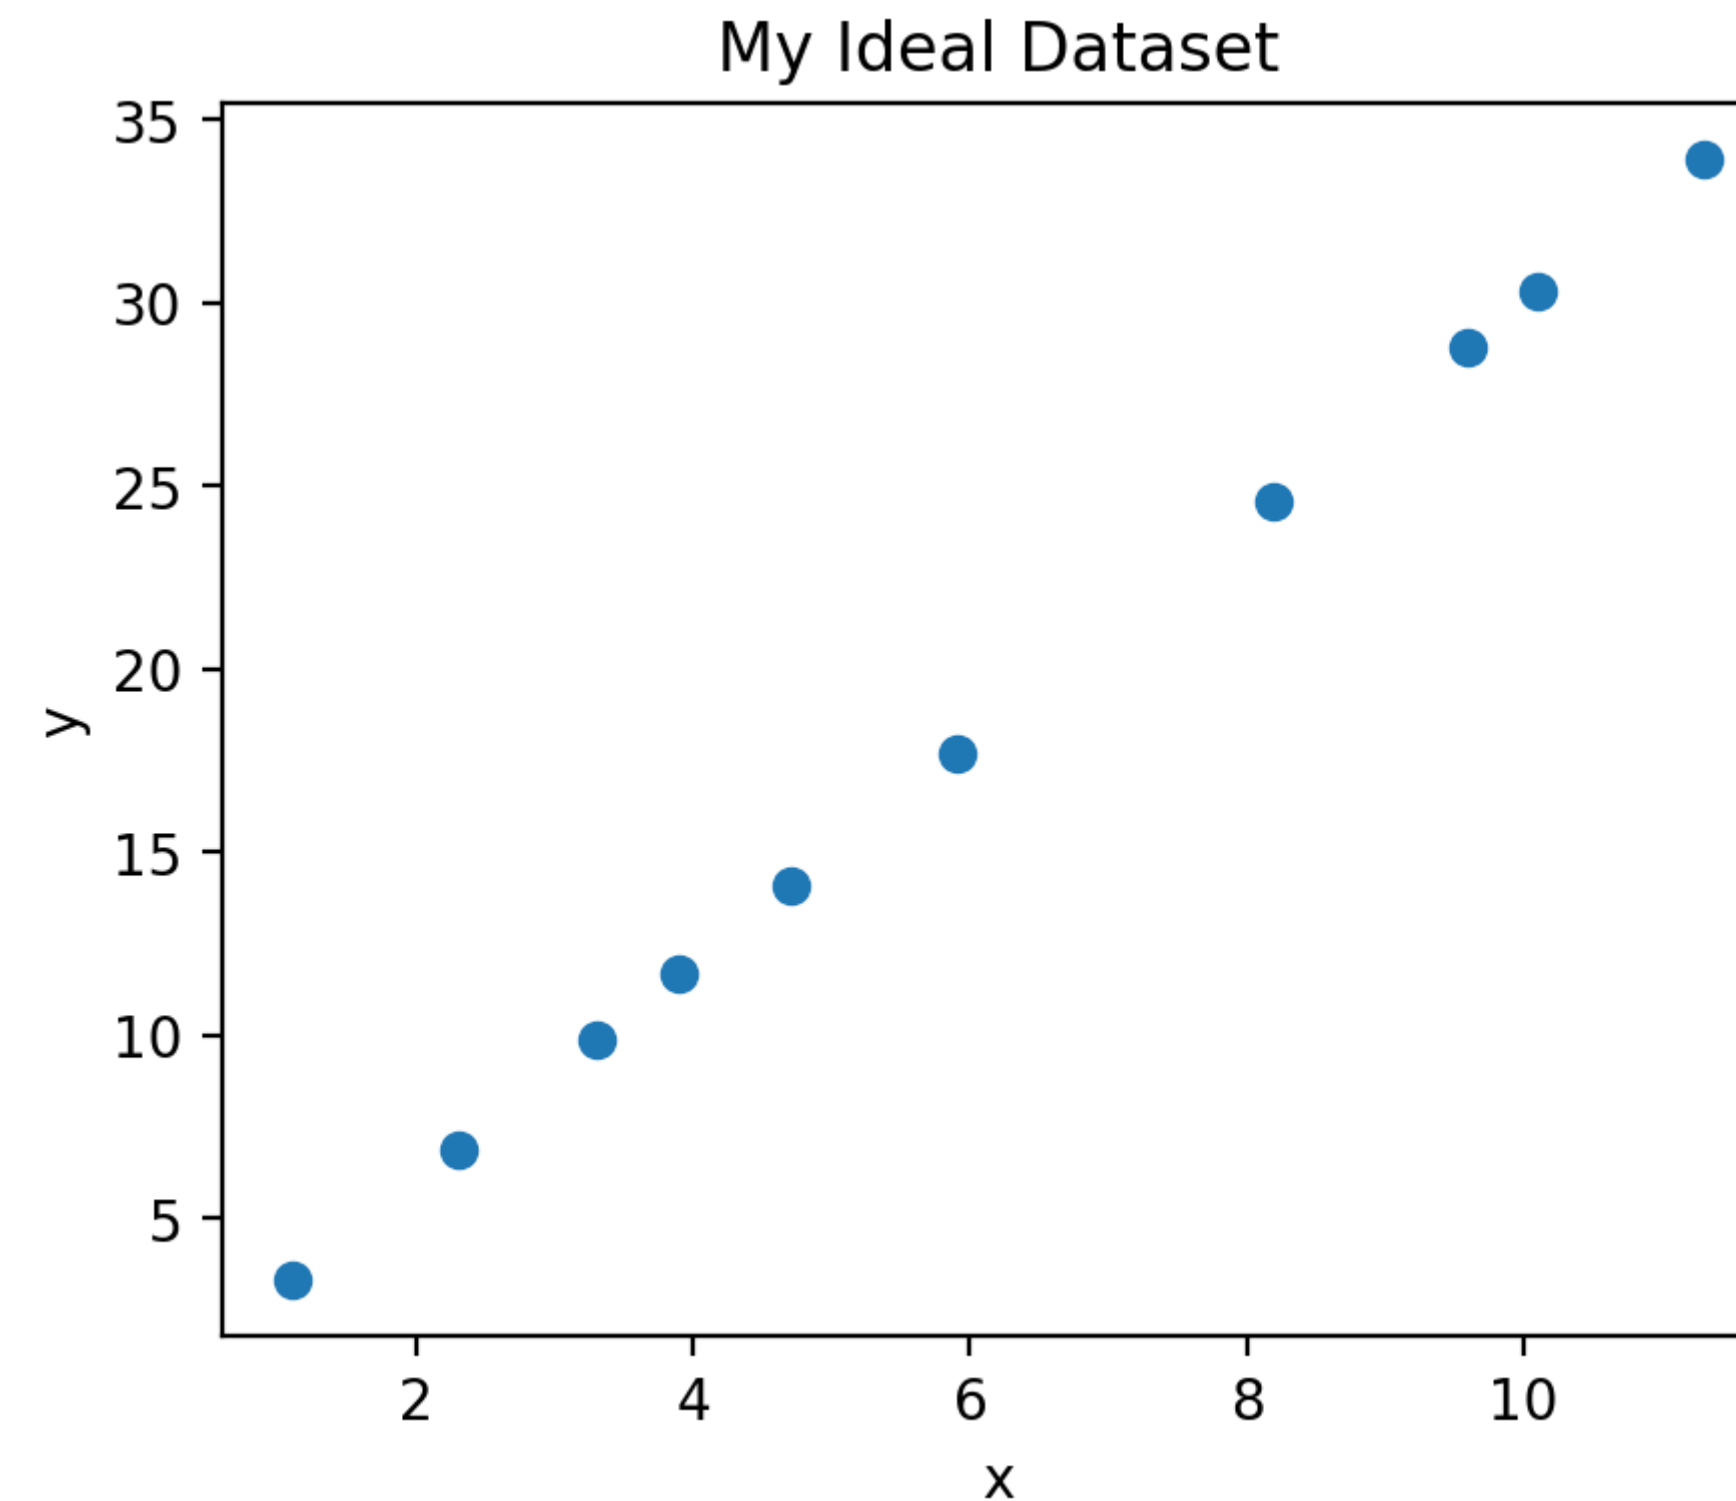
Bayesian Neural Networks

Practical Deep Learning for Science
20 June, 2024

Let's say we are doing an experiment

X	Y
1.1	3.3
2.3	6.9
3.3	9.9
3.9	11.7
9.6	28.8
10.1	30.3
11.3	33.9
4.7	14.1
8.2	24.6
5.9	17.7

What is Y at $x = 7$?

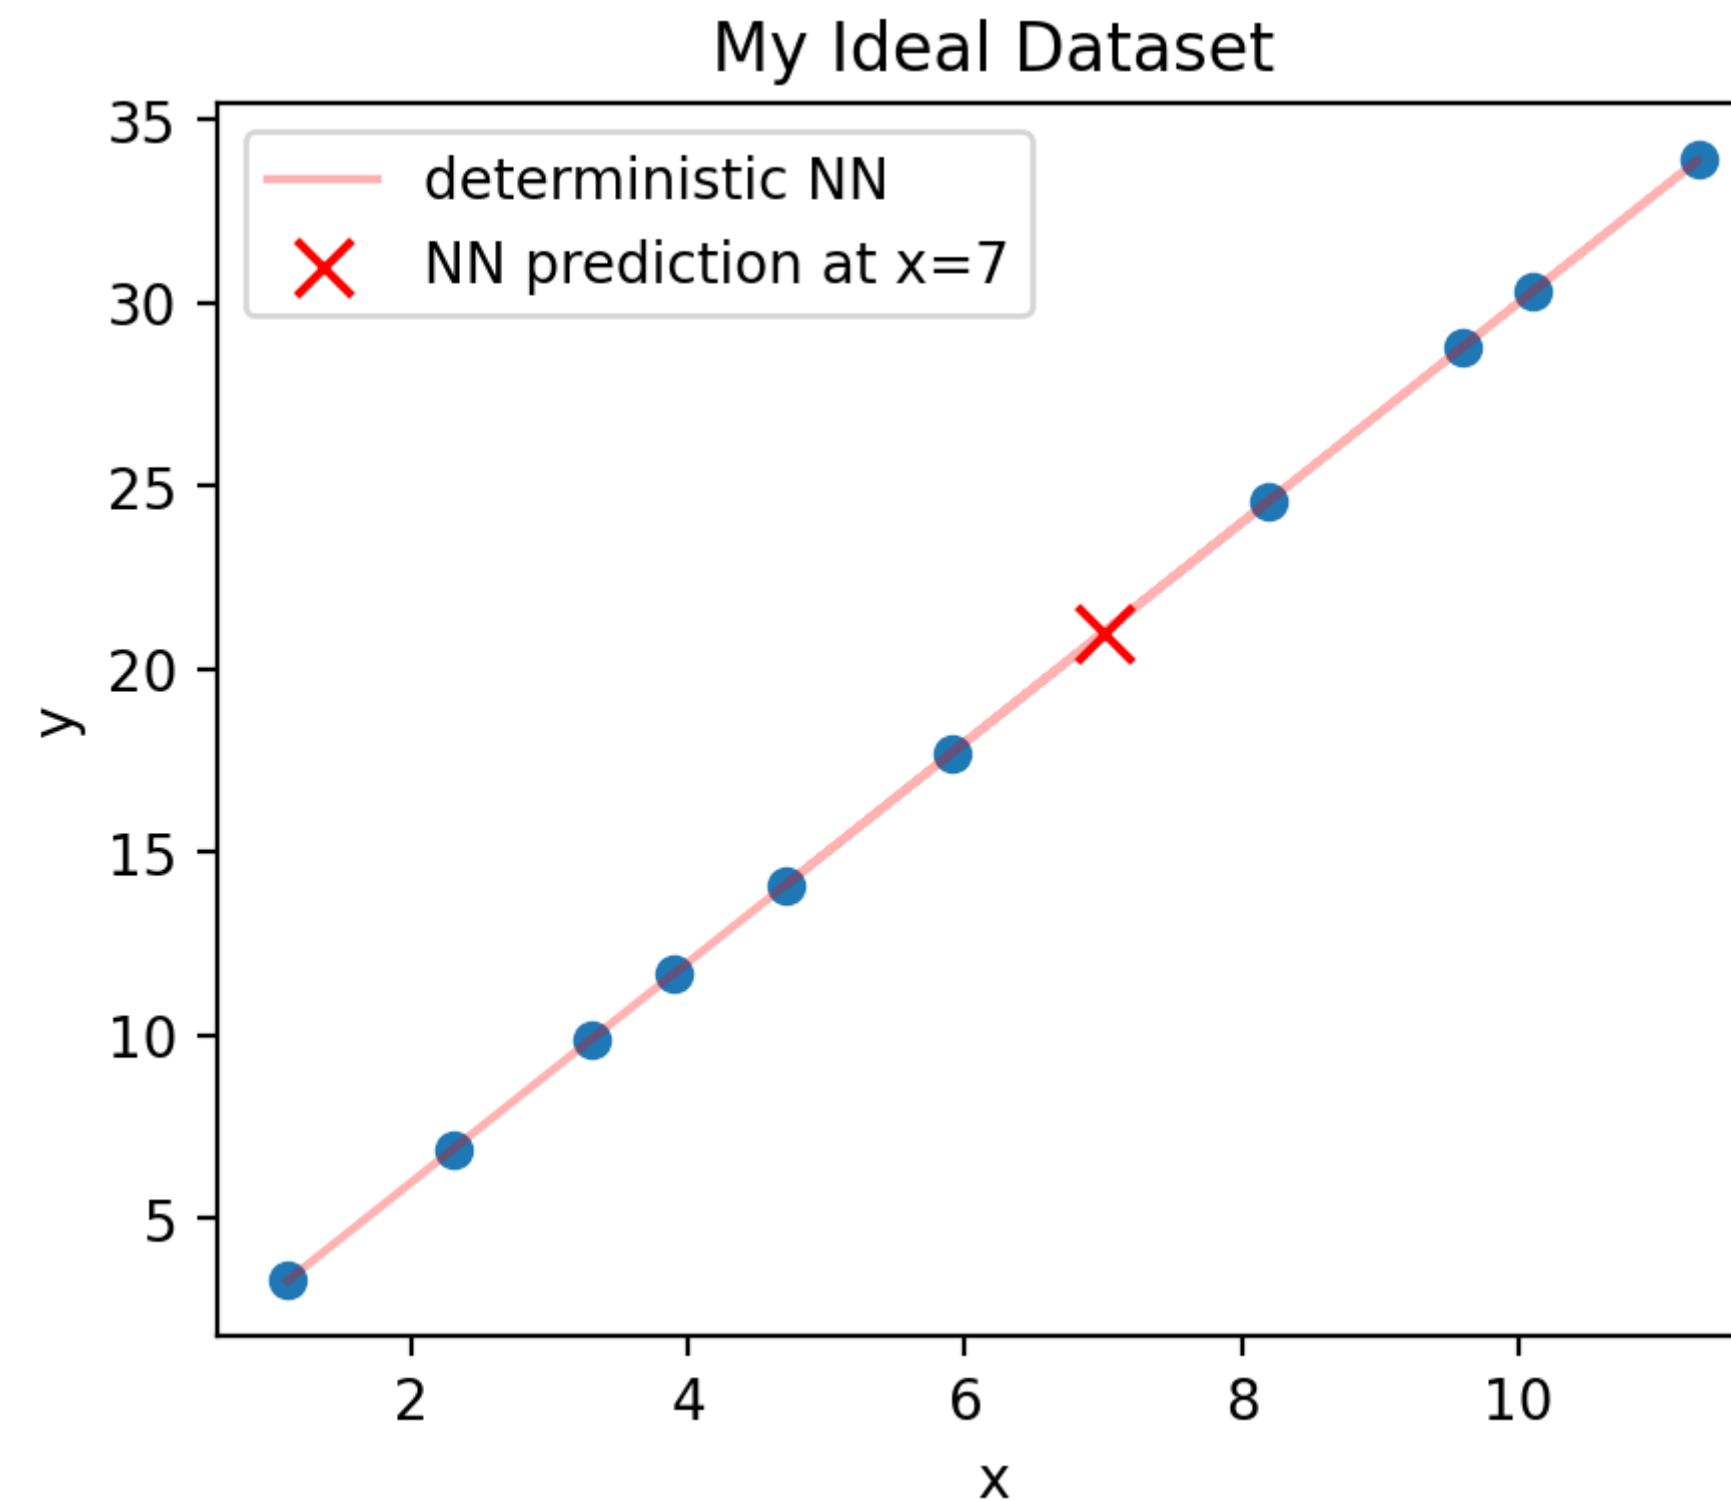


- We can still train an NN with only one weight to do the same

Let's say we are doing an experiment

X	Y
1.1	3.3
2.3	6.9
3.3	9.9
3.9	11.7
9.6	28.8
10.1	30.3
11.3	33.9
4.7	14.1
8.2	24.6
5.9	17.7


What is Y at $x = 7$?




- We can still train an NN with only one weight to do the same

Let's say we are doing an experiment

Ideal situation




What we see




If we train a network

We will get



What we want



How do we quantify the uncertainty with NN

Aleatoric uncertainty

- Noise in the data
- More data won't help


vs

Epistemic uncertainty


- Uncertainty in the model parameters
- More data will help

- We need to modify the network to accommodate these two uncertainties

Model uncertainty



Model weight before




Gaussian (μ, σ)

Model weight after

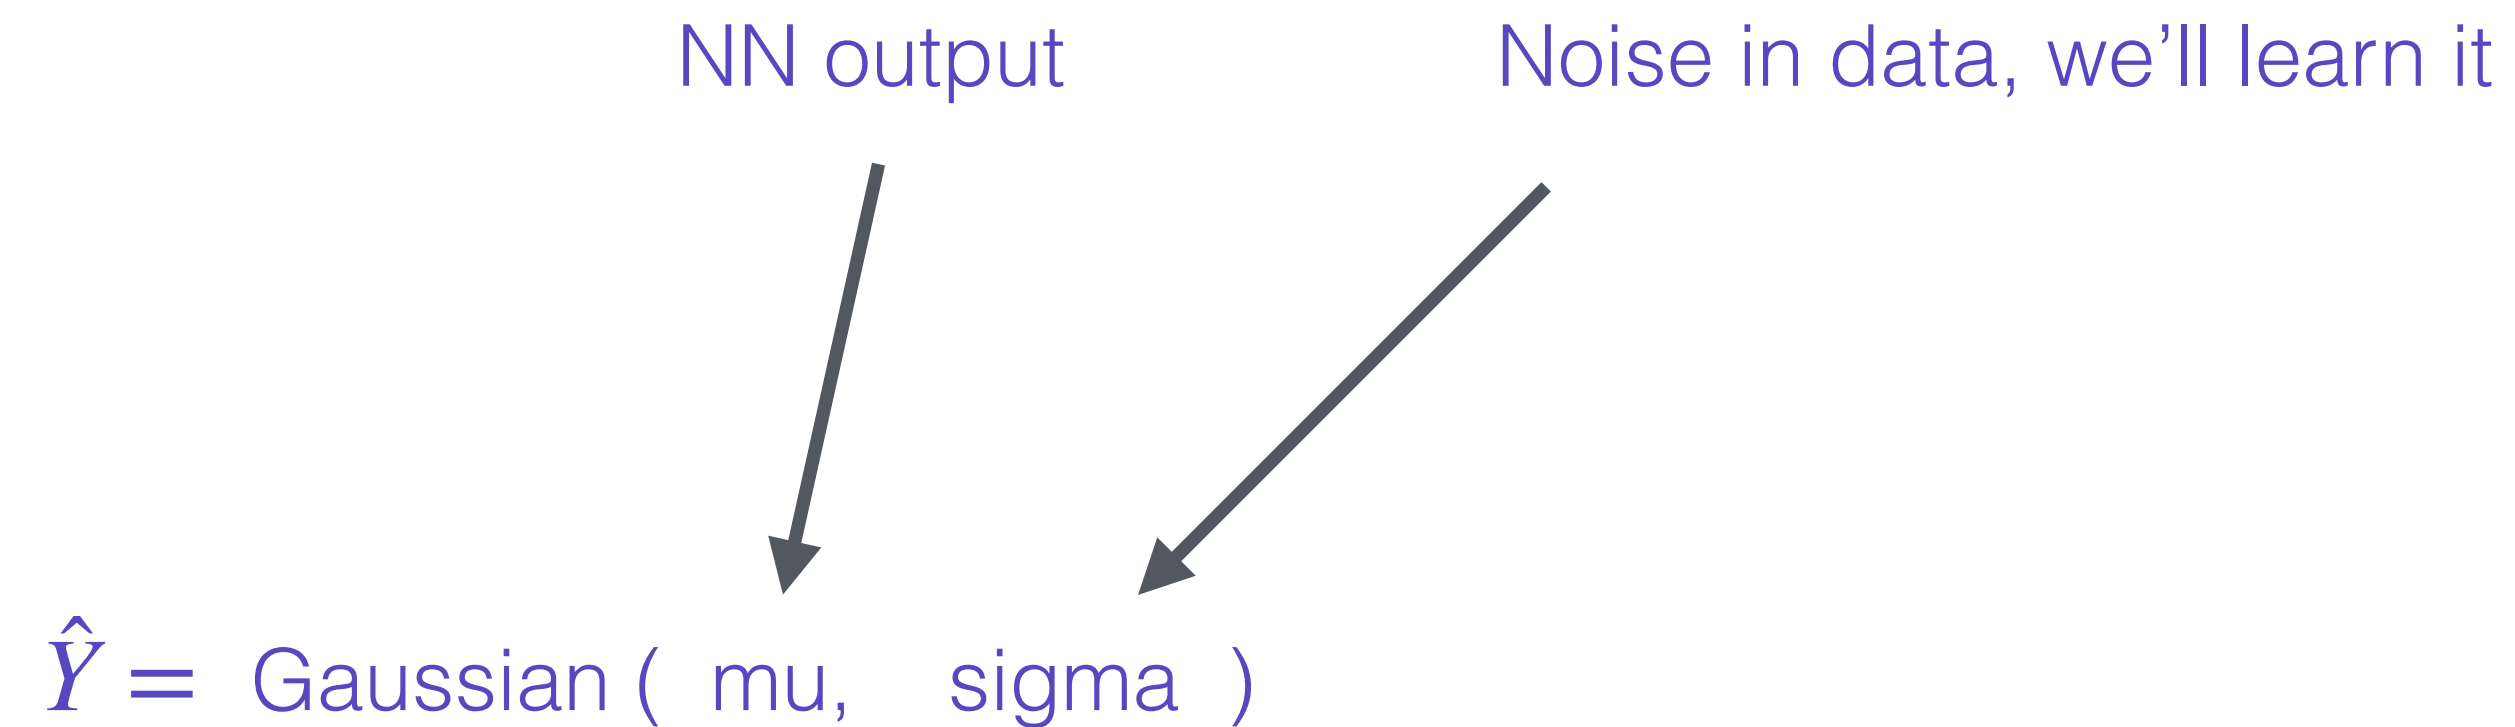
- ◆ Each weight, instead of being a number, will be two numbers - μ and σ
- ◆ During forward pass, we sample from the Gaussian (μ, σ)

Model uncertainty



How about the uncertainty from the data

- Let's add another Gaussian into the picture
- Mean = output of the NN (a gaussian)
- Sigma = another Gaussian we'll learn


$$\hat{Y} = \text{Gaussian}(\mu, \sigma)$$

We are drawing a Gaussian around the network prediction,
and we'll learn the width of this Gaussian

‘Bayes’ian NN

Li'l bit of extra maths :)

Bayesian networks

♦ We have

→ θ network parameters

→ D data

$$p(x) = \text{distribution of } x$$

♦ We want to get the right network parameters, given the data (through training)

→ $p(\theta | D)$, we don't know how to compute it directly

→ Bayes' theorem gives us

$$p(\theta | D) = \frac{p(D | \theta) p(\theta)}{p(D)}$$

→ $p(y | x, \theta)$ (model output distribution given input and NN)

→ Prior (initial distribution of the NN weights; mostly Gaussian)

→ Distribution of data

$p(D)$ is expensive!

$$\diamond p(D) = \int p(D|\theta) p(\theta) d\theta$$

- ➔ Need to compute over all possible values of θ
- ➔ Not possible!

♦ So, we approximate

- ➔ Variational Inference (VI)

$p(x)$ = distribution of x

Approximate the NN weight distribution $p(\theta|D)$

- ◆ Let's assume $p(\theta|D)$ follows some distribution (say, Gaussian)
 - ➔ We try to approximate its parameters
 - ➔ $p(\theta|\phi)$ (how my NN weights are distributed, given the parameters)
 - ➔ ϕ will be μ and σ for our Gaussian p
- ◆ We want to make approximate posterior, $p(\theta|\phi)$ close to the true posterior, $p(\theta|D)$
 - ➔ So, ideally a KLD b/w the two, but we of course don't know the true posterior
- ◆ Turns out, this is equivalent to maximizing ELBO (Evidence Lower bound)

$$ELBO = \mathbb{E} \left[\log p(D|\theta) \right] - KL \left(p(\theta|\phi) \parallel p(\theta) \right)$$

```
svi = SVI(model, guide, adam, loss=Trace_ELBO())
```