

# Bayesian Neural Networks

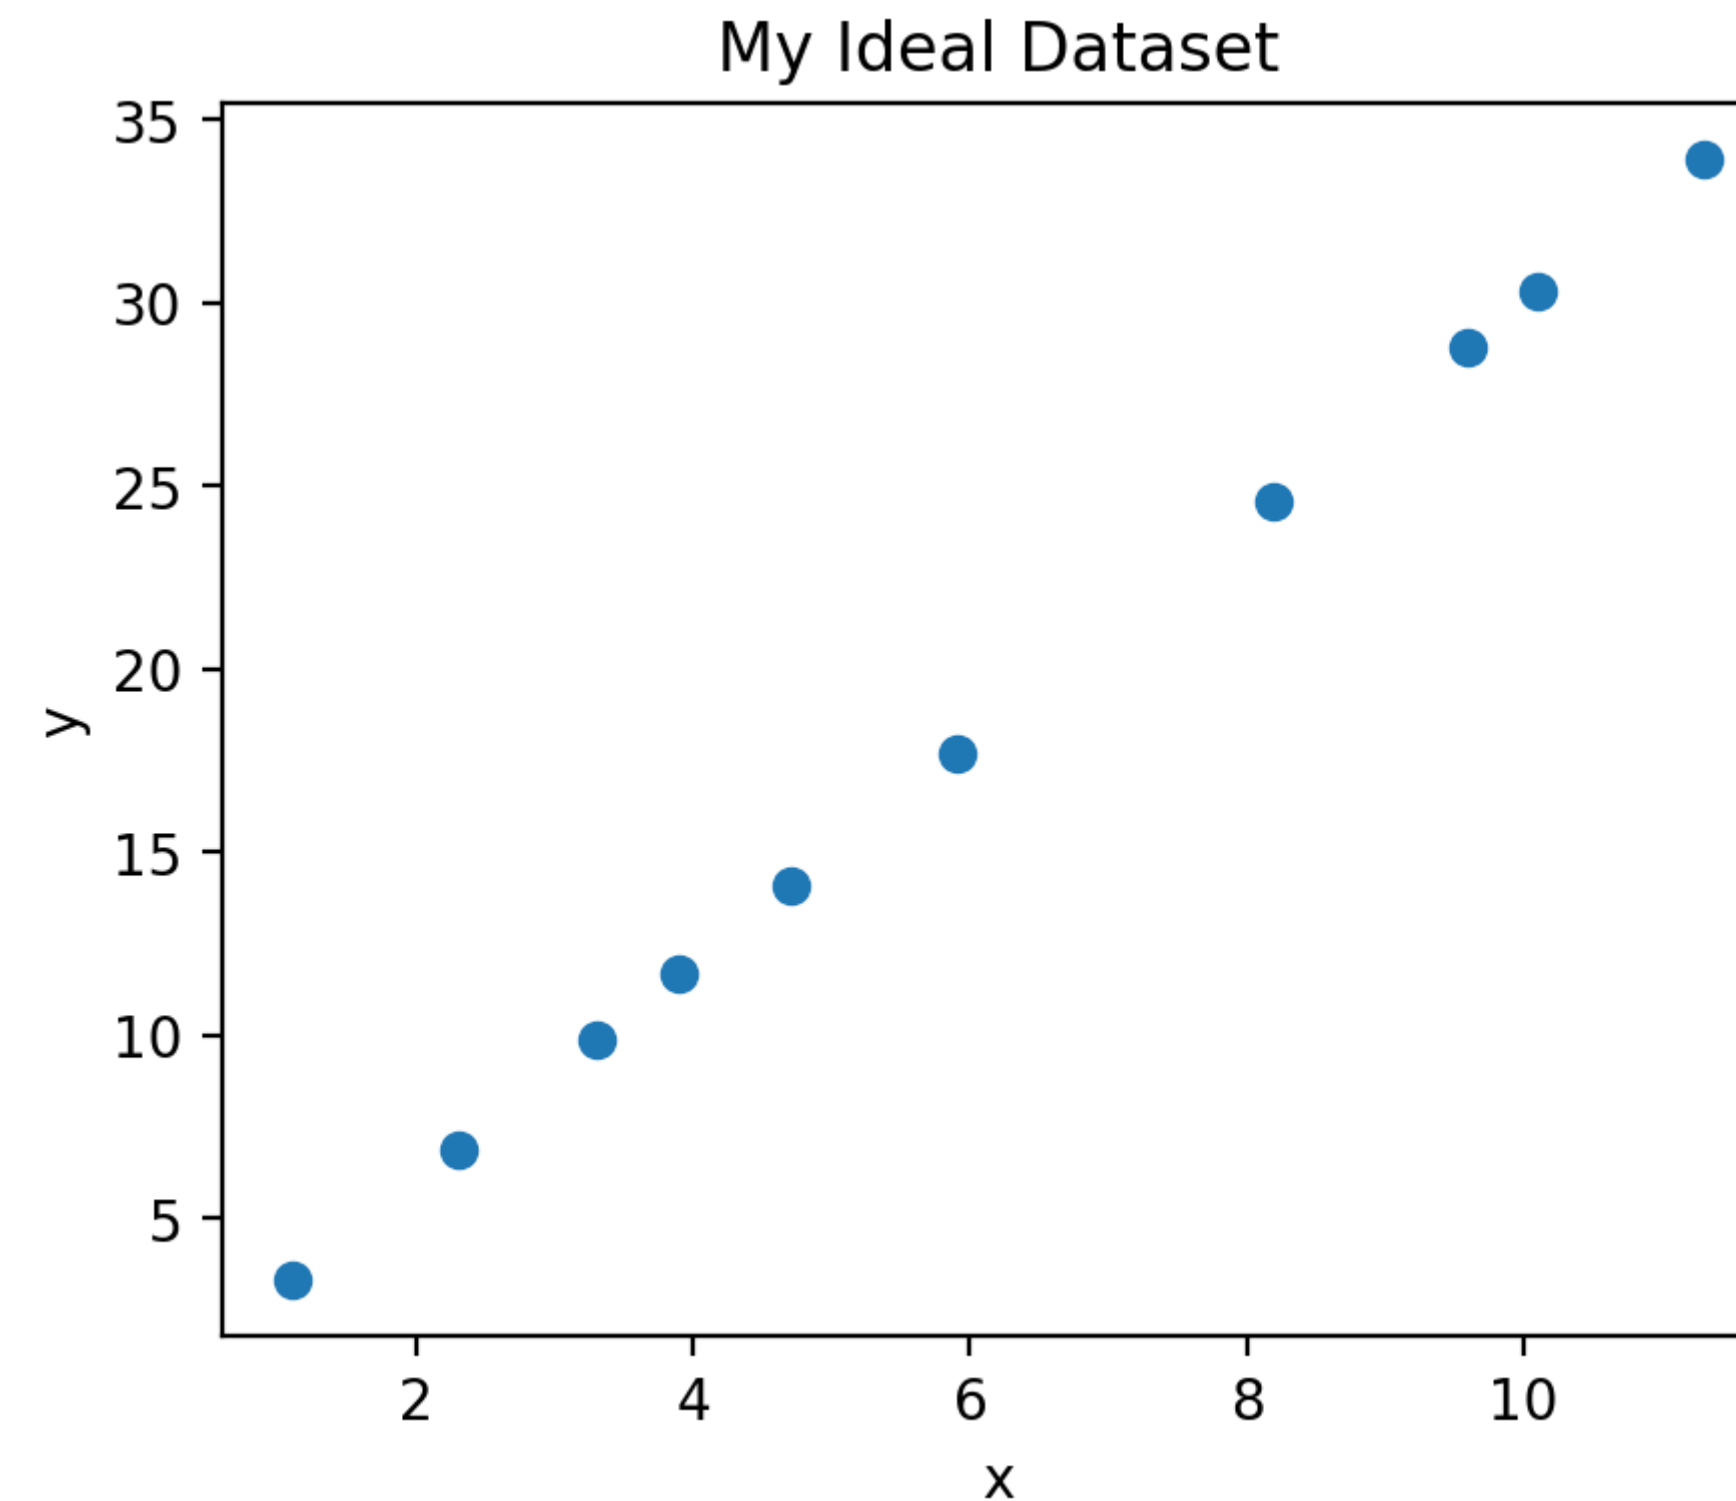
---

**Practical Deep Learning for Science**  
**20 June, 2024**

# Let's say we are doing an experiment

X	Y
1.1	3.3
2.3	6.9
3.3	9.9
3.9	11.7
9.6	28.8
10.1	30.3
11.3	33.9
4.7	14.1
8.2	24.6
5.9	17.7

What is Y at  $x = 7$ ?

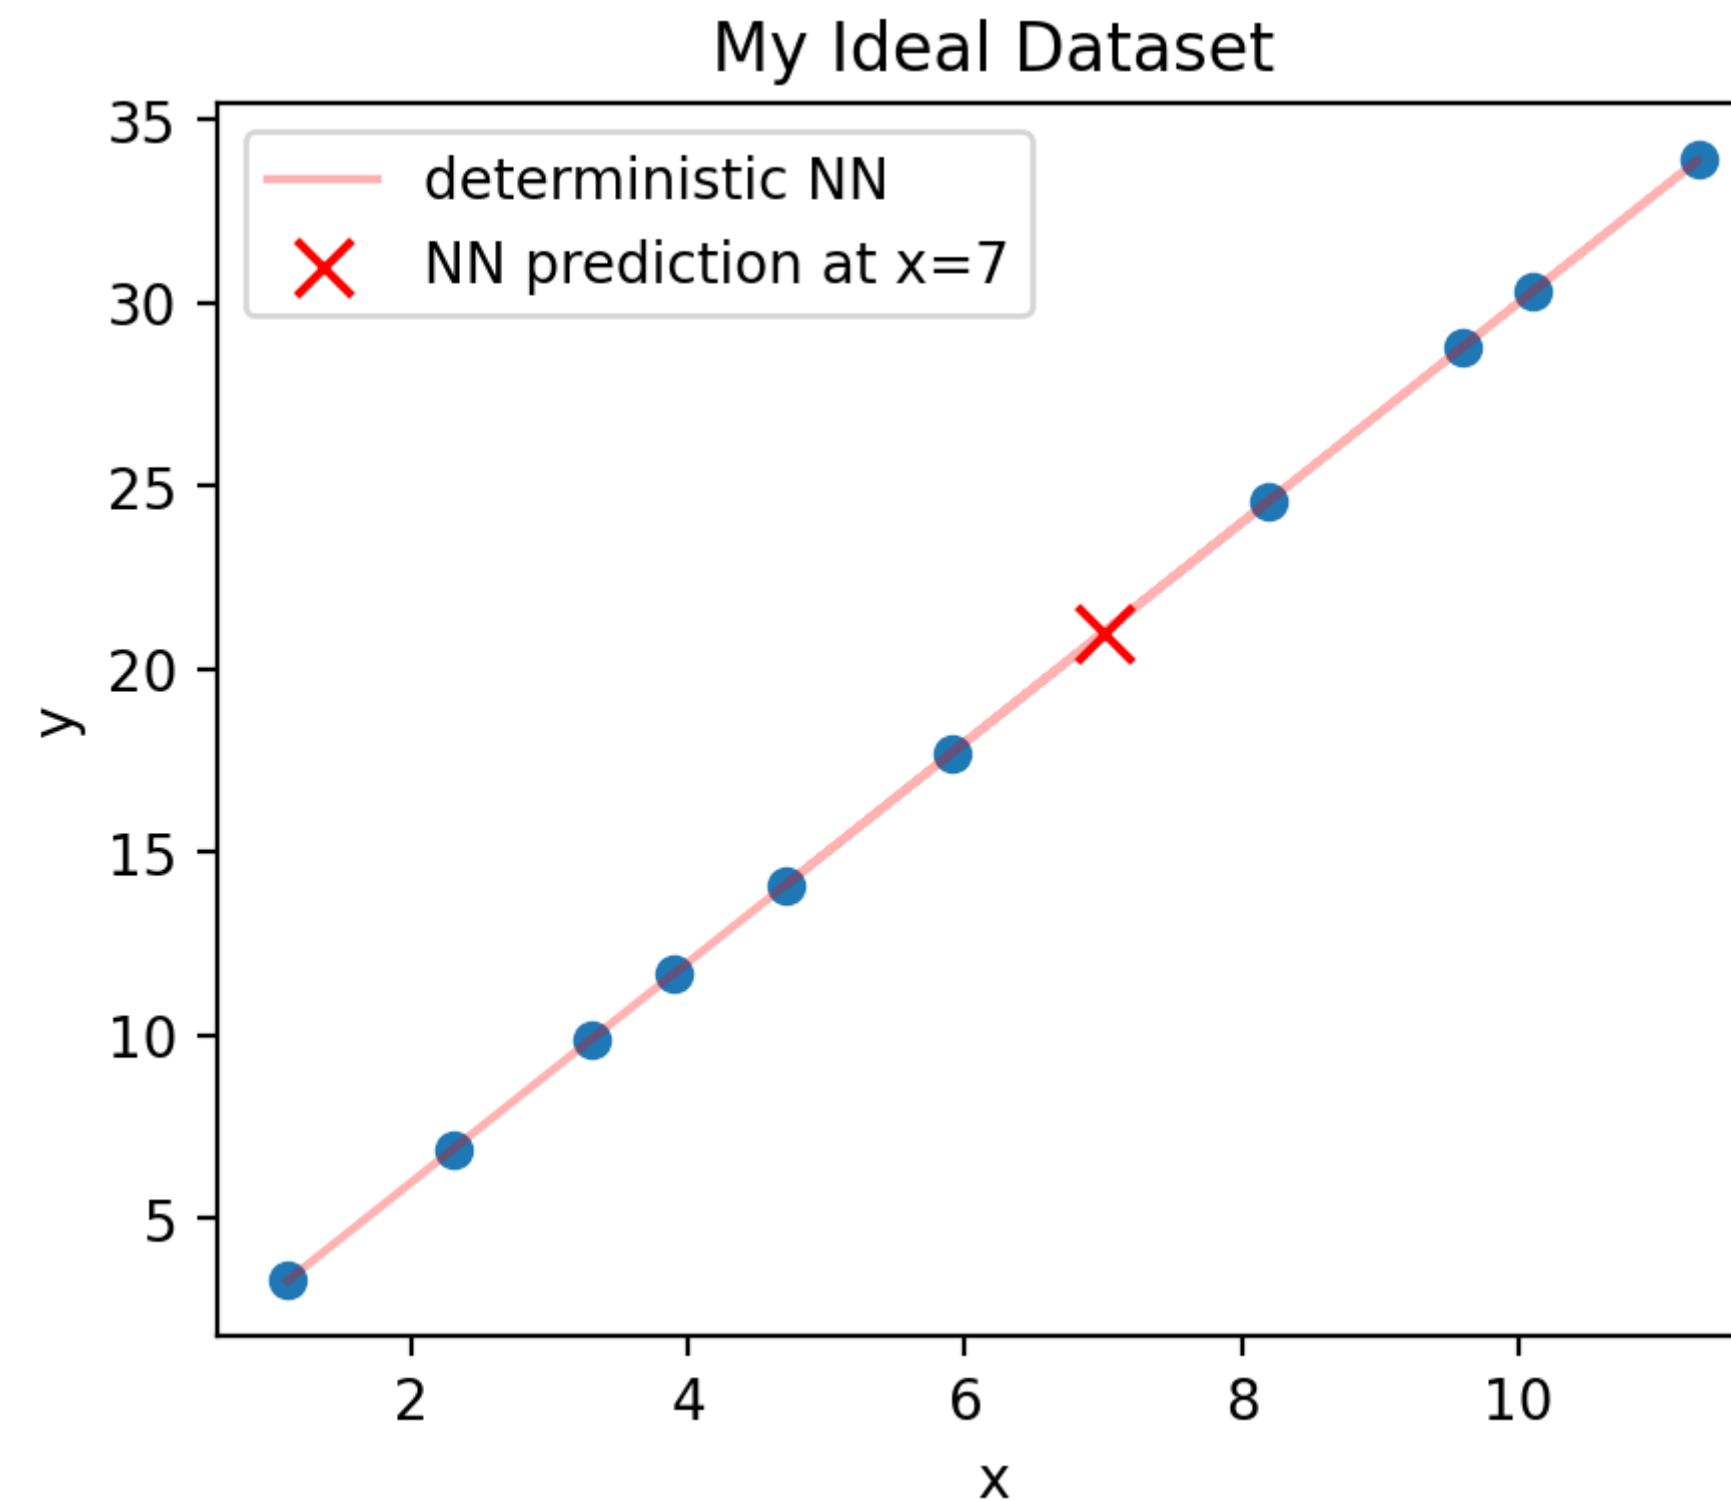


- We can still train an NN with only one weight to do the same

# Let's say we are doing an experiment

X	Y
1.1	3.3
2.3	6.9
3.3	9.9
3.9	11.7
9.6	28.8
10.1	30.3
11.3	33.9
4.7	14.1
8.2	24.6
5.9	17.7


What is Y at  $x = 7$ ?




- We can still train an NN with only one weight to do the same

# Let's say we are doing an experiment

## Ideal situation




## What we see

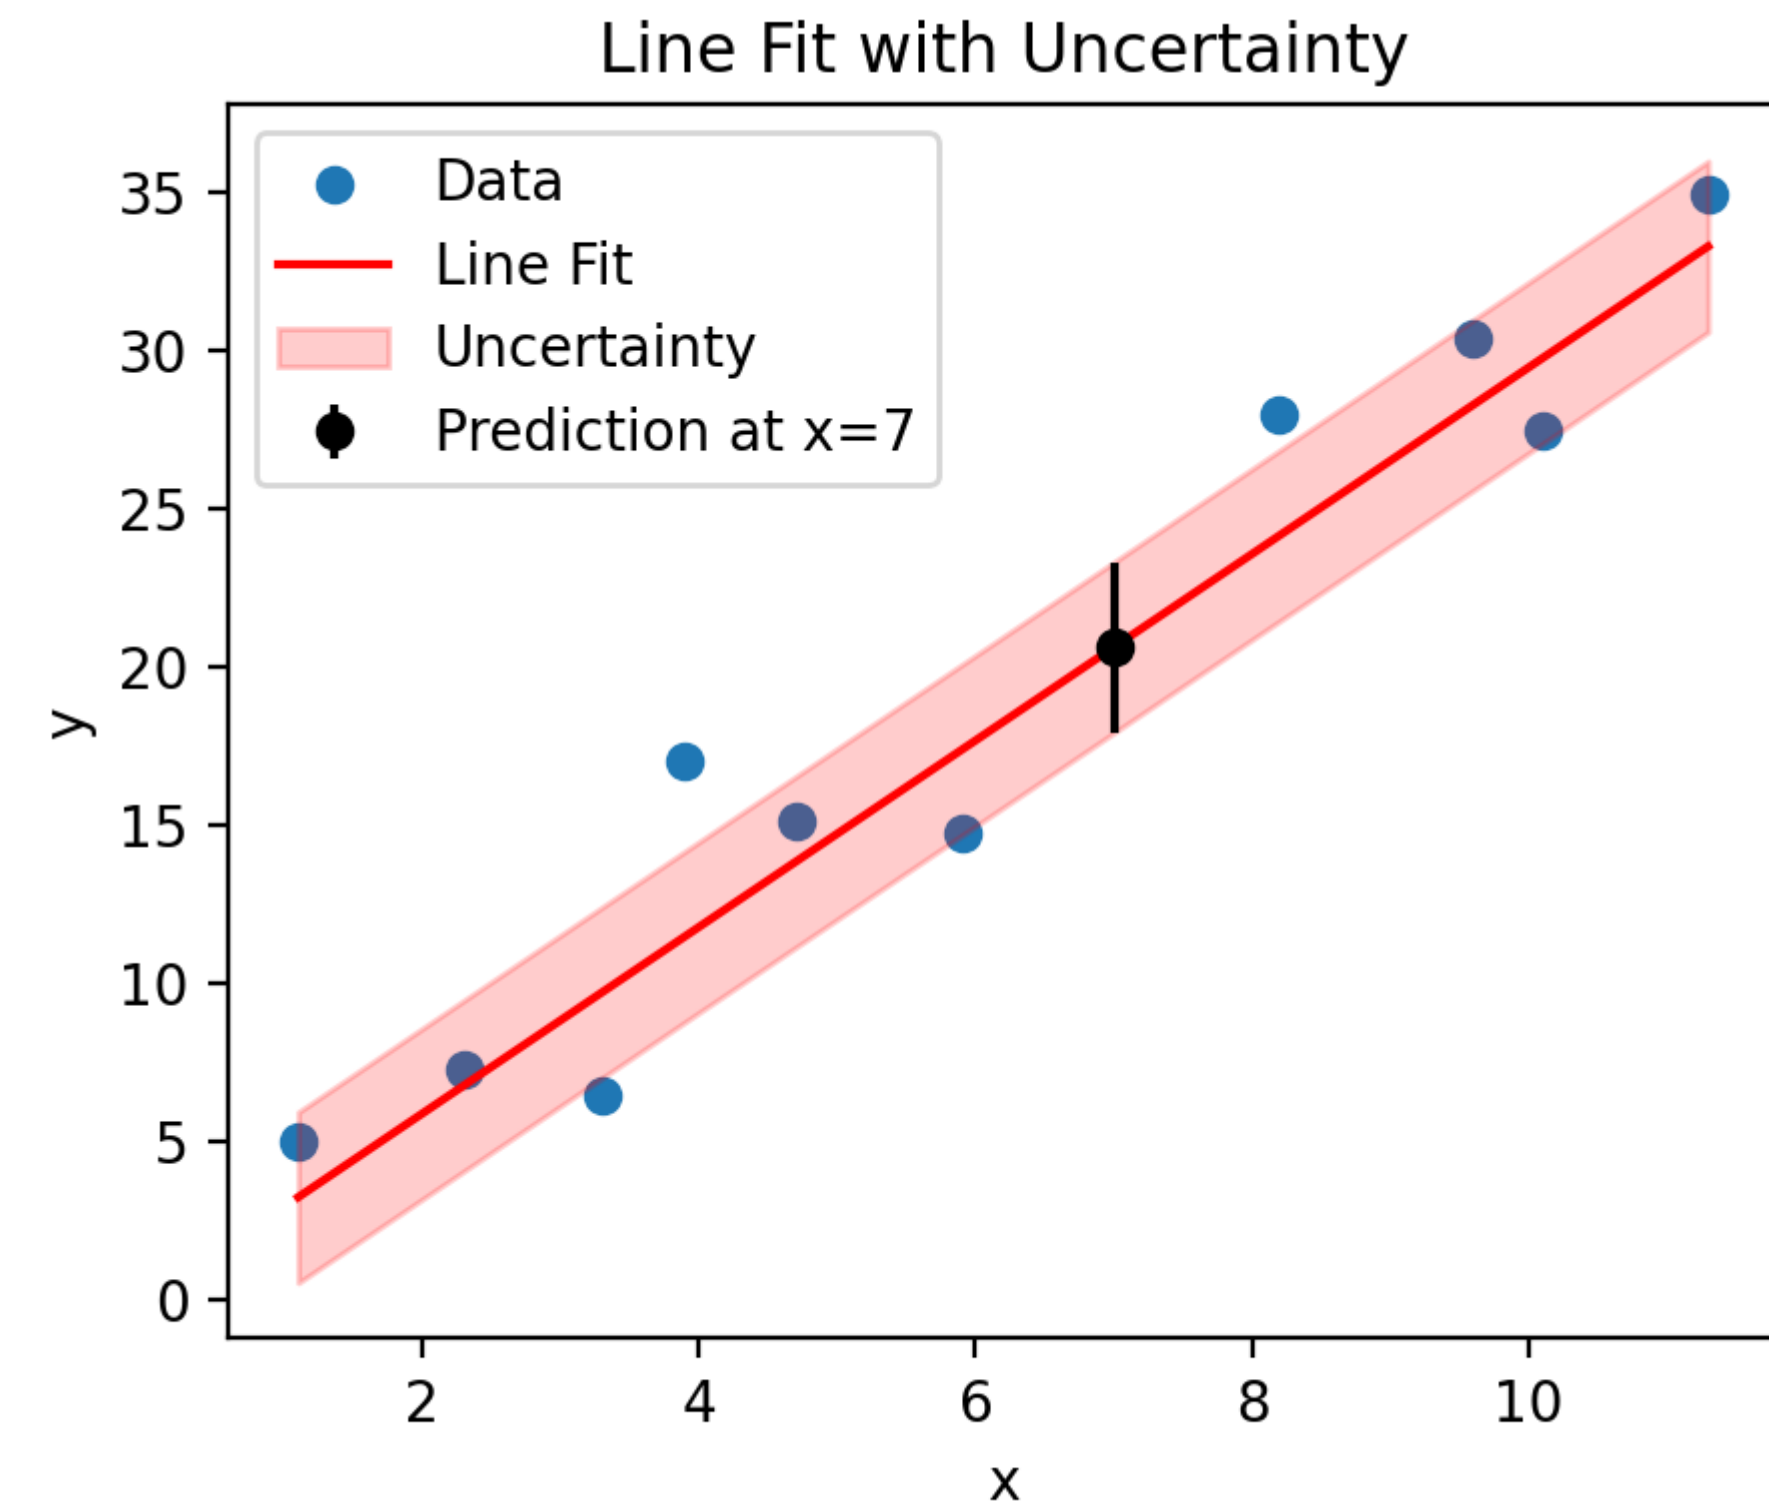


# If we train a network

We will get



What we want



# How do we quantify the uncertainty with NN

## Aleatoric uncertainty

- Noise in the data
- More data won't help


**vs**

## Epistemic uncertainty


- Uncertainty in the model parameters
- More data will help

- We need to modify the network to accommodate these two uncertainties

# Model uncertainty



Model weight before





Gaussian ( $\mu, \sigma$ )

Model weight after

- ◆ Each weight, instead of being a number, will be two numbers -  $\mu$  and  $\sigma$
- ◆ During forward pass, we sample from the Gaussian ( $\mu, \sigma$ )


# Model uncertainty





# How about the uncertainty from the data

- Let's add another Gaussian into the picture
- Mean = output of the NN (a gaussian)
- Sigma = another Gaussian we'll learn



We are drawing a Gaussian around the network prediction,  
and we'll learn the width of this Gaussian

# ‘Bayes’ian NN

---

Li'l bit of extra maths :)

# Bayesian networks

♦ We have

→  $\theta$  network parameters

→  $D$  data

$$p(x) = \text{distribution of } x$$

♦ We want to get the right network parameters, given the data (through training)

→  $p(\theta | D)$ , we don't know how to compute it directly

→ Bayes' theorem gives us

$$p(\theta | D) = \frac{p(D | \theta) p(\theta)}{p(D)}$$

→  $p(y | x, \theta)$  (model output distribution given input and NN)

→ Prior (initial distribution of the NN weights; mostly Gaussian)

→ Distribution of data

# $p(D)$ is expensive!

$$\blacklozenge \quad p(D) = \int p(D|\theta) p(\theta) d\theta$$

- ➔ Need to compute over all possible values of  $\theta$
- ➔ Not possible!

✦ So, we approximate

- ➔ Variational Inference (VI)

$p(x)$  = distribution of  $x$

# Approximate the NN weight distribution $p(\theta|D)$

- ◆ Let's assume  $p(\theta|D)$  follows some distribution (say, Gaussian)
  - ➔ We try to approximate its parameters
  - ➔  $p(\theta|\phi)$  (how my NN weights are distributed, given the parameters)
    - ➔  $\phi$  will be  $\mu$  and  $\sigma$  for our Gaussian  $p$
- ◆ We want to make approximate posterior,  $p(\theta|\phi)$  close to the true posterior,  $p(\theta|D)$ 
  - ➔ So, ideally a KLD b/w the two, but we of course don't know the true posterior
- ◆ Turns out, this is equivalent to maximizing ELBO (Evidence Lower bound)

$$ELBO = \mathbb{E} \left[ \log p(D|\theta) \right] - KL \left( p(\theta|\phi) || p(\theta) \right)$$

```
svi = SVI(model, guide, adam, loss=Trace_ELBO())
```